

# **PARALLEL GRAPH COLORING ALGORITHMS**


Margomenos Spyridon  
Denzler Alain  
Moesch Philippe

# Graph coloring

- Given  $G = (V, E) \rightarrow$  for all  $(u, v)$  in  $E$  :  
color  $(u) \neq$  color  $(v)$
- Goal is to use as few colors as possible
- Problem is NP-Complete  $\rightarrow$  Faster algorithms approximate the optimal solution

# Boman's graph coloring algorithm

- Iterative distributed-memory graph coloring algorithm
- Vertices are split among  $p$  processors



- Tentative coloring phase :

- Each processors color its vertices, should be valid according to the previous global color allocation
- ! neighbors on different processors might be assigned the same color → invalid allocation
- communication regarding boundary vertices required (batched in supersteps)

- Conflict resolution phase :

- For each conflict, one of the two vertices gets marked as to be recolored in the next tentative coloring phase
- Eventually there will be no more conflicts
- No communication to be made in this phase

# Tunable Parameters

- Superstep size

Tradeoff between large superstep inducing less communication overhead but more conflicts and smaller superstep with more frequent communication

- Communication style


Communication can either be synchronous or asynchronous (obvious latency advantage, but more conflicts with asynchronous)

# Color assignment algorithms

- First Fit

Each processors chooses the smallest admissible color in  $[1, N]$ , where  $N$  is the largest color used. If no color admissible choose  $N+1$

- Staggered First Fit (N initial estimate)



If no admissible look from 1 to  $\text{floor}(i \cdot N/p)$

If still no admissible choose smallest permissible color  $> N$

Search first from  $\text{ceiling}(i \cdot N/p)$  to  $\text{ceiling}((i+1) \cdot N/p)$  and then

From  $\text{ceiling}(i \cdot N/p)$  to 0

- Modified Staggered First Fit

# Matrix Storage Approach

Slightly modified CSR format with no values array

	0	1	2	3	4
0	1		1		1
1		1		1	
2		1	1		
3	1		1	1	
4		1		1	

	0	1	2	3	4	5						
rowptr	0	3	5	7	10	12						
	0	1	2	3	4	5	6	7	8	9	10	11
colind	0	2	4	1	3	1	2	0	2	3	1	3



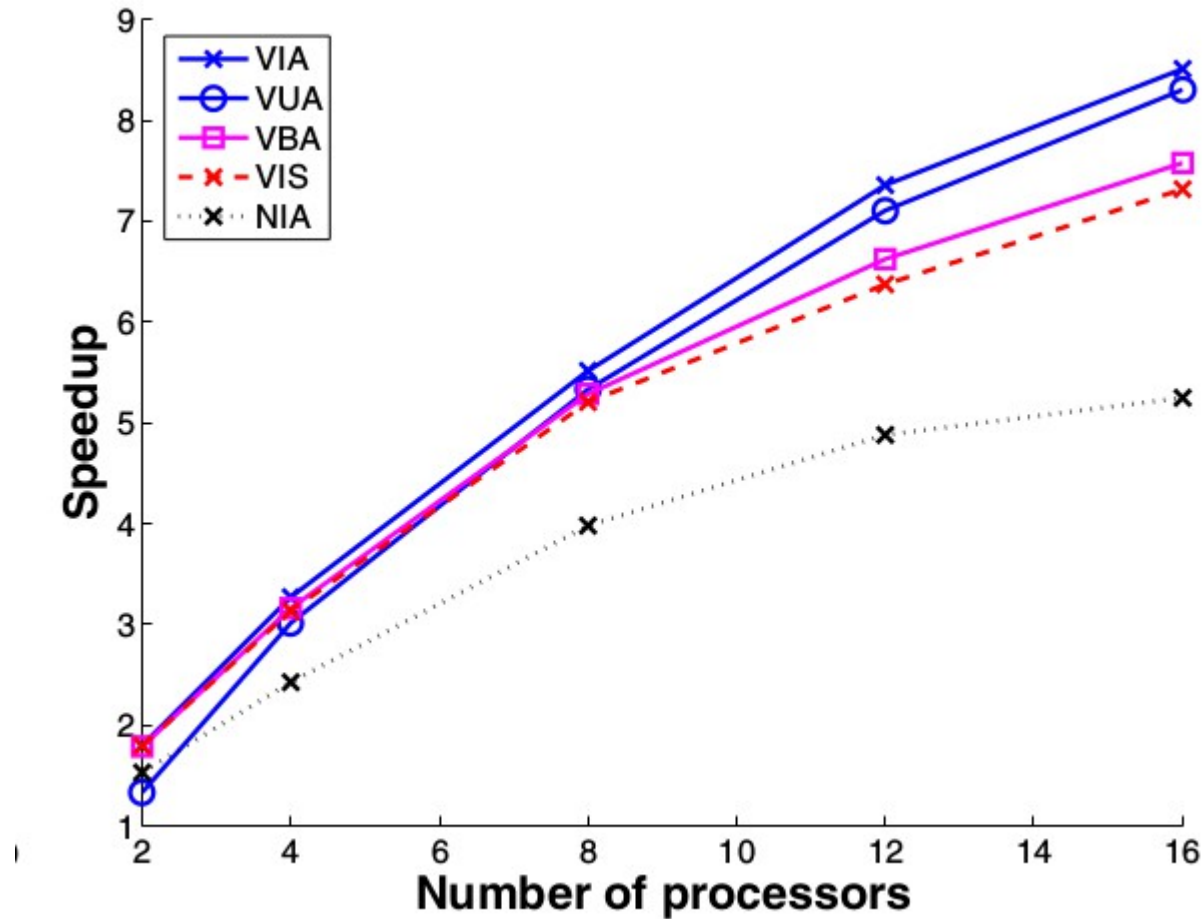
# Implementation Details

- C/C++ & MPI
- Focus on regular and predictable array accesses
- Compact data structures
- Addition of modified staggered first fit (“conflict-avoiding color fitting”)

# Previous work

- The original paper (2005) proposed the following :
  - First-fit coloring assignment
  - 100 superstep size
  - assumes graphs are well-partitioned

# Result




16 nodes PC-cluster with dual 900 MHz Intel Itanium 2 CPUs & 4 GB memory

# Test Graphs

- RMAT random graphs
- 3 size levels: 1M/10M/30M vertices
- 3 density levels: 1:3/7/15 Vertices:Edges
- 10 graphs for each level-density combination
- Assumption: graph is distributed and knows about topology
- 90 graphs vs 7 configurations vs superstep sizes – unfortunately did not finish in time


Absolute Runtime 1M7, Superstep 500




# Test Graphs

- Instead: real world graphs run on Euler
- Californian Road Network: 1.9M/2.8M
- Google Web Graph: 900K/8.6M
- Livejournal Social Network: 4M/69M
- Orkut Social Network: 3M/234M

Strong Scaling CA-Roadnet (1.9M/2.8M), Superstep 500

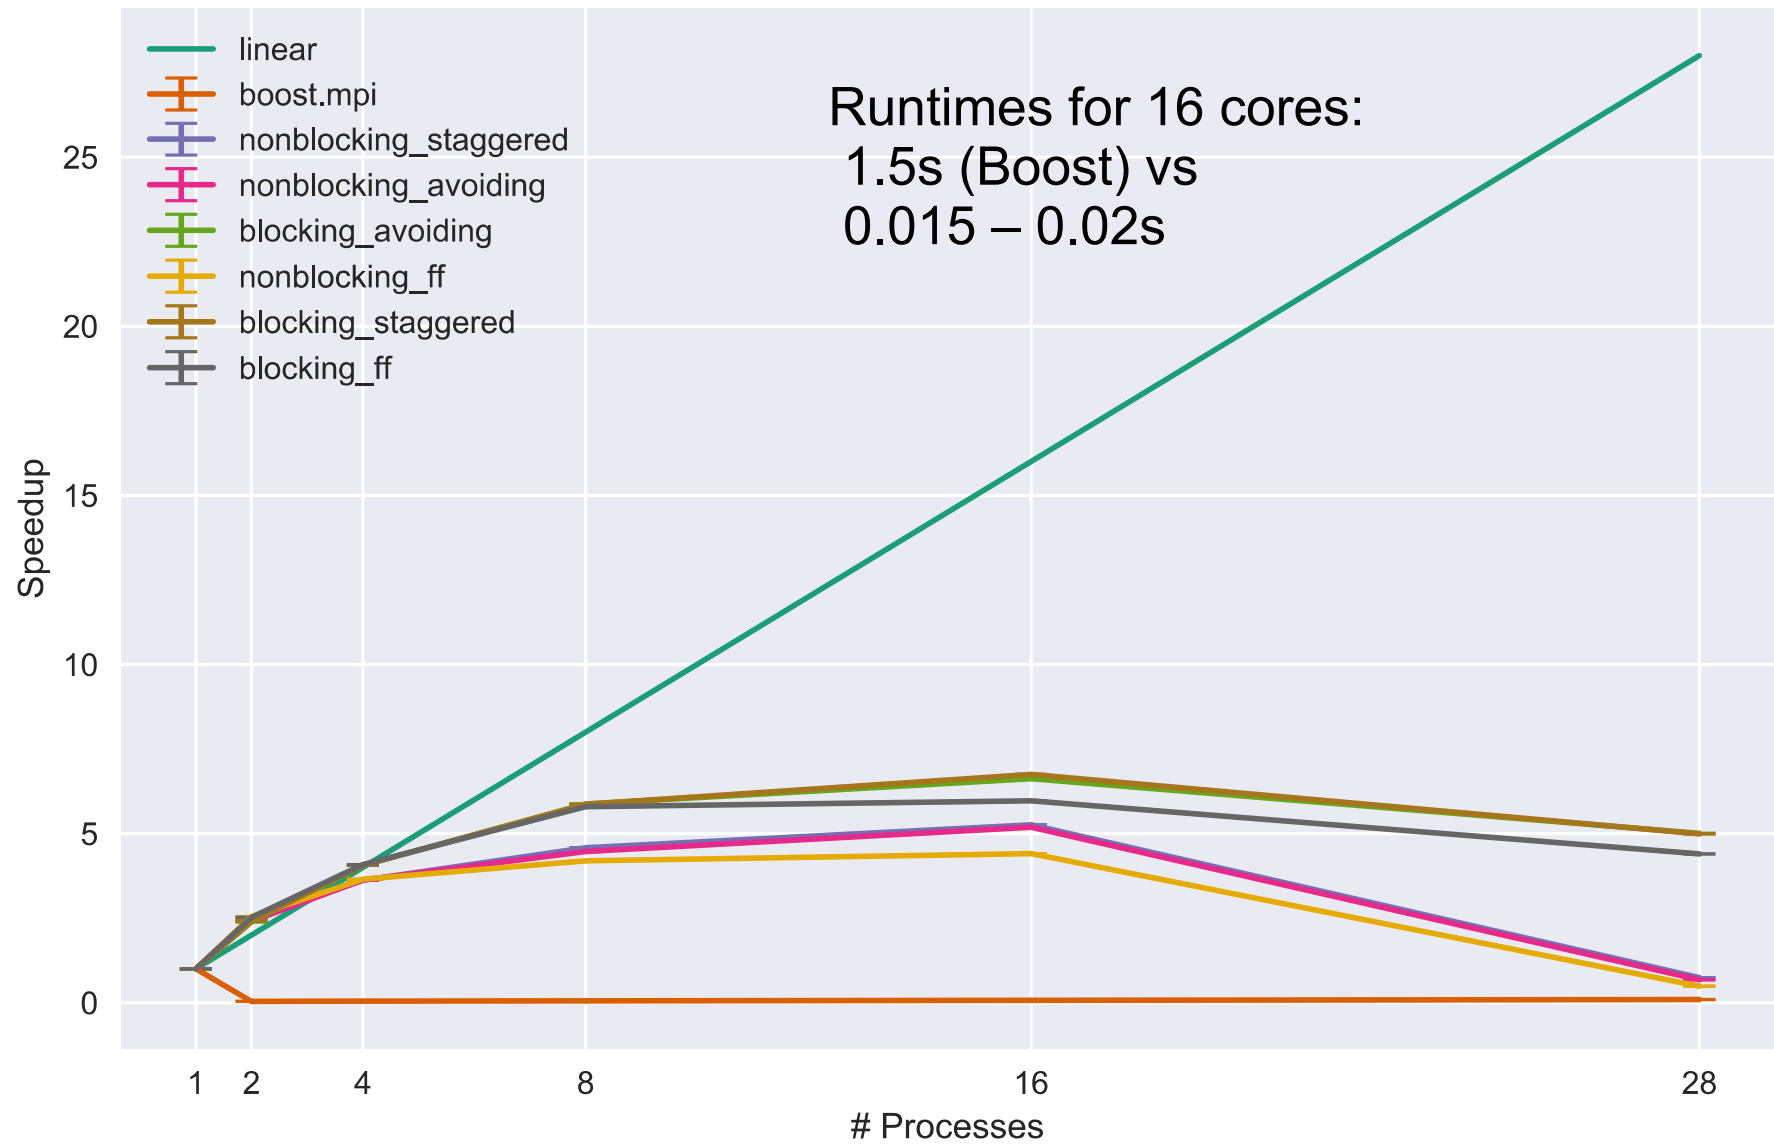


Absolute Runtime CA-Roadnet (1.9M/2.8M), Superstep 500







Strong Scaling Google (900K/8.6M), Superstep 500




Strong Scaling LiveJournal (4M/69M), Superstep 500




Strong Scaling Orkut (3M/234M), Superstep 500



Absolute Runtime Orkut (3M/234M), Superstep 500



Normalized Number of Colors, Orkut (3M/234M), Superstep 500



# Conclusion

- Our implementation is really fast
- Better scaling on denser/larger graphs
- Blocking scales better than nonblocking
- Tradeoff complexity – speed – number of colors in color assignment